



端到端效率与规模扩展 (冻结 V2; 每档 5 次)


A. 五档规模的端到端耗时




B. 100% 规模的阶段中位耗时



C. warm-process 五点拟合残差



D. 进程峰值内存



warm 不含模型包加载; 含加载口径纳入每次真实模型包加载, 但未清空 OS 文件缓存。缓存输出与在线 LLM 均不计入; 拟合仅为五点描述。